


a Correct runs for three for every task, model and condition



b Execution errors by type and channel (all seven types)



c Final correctness by execution errors, per benchmark

